# Deduce the World: How Far Can Reasoning Go?

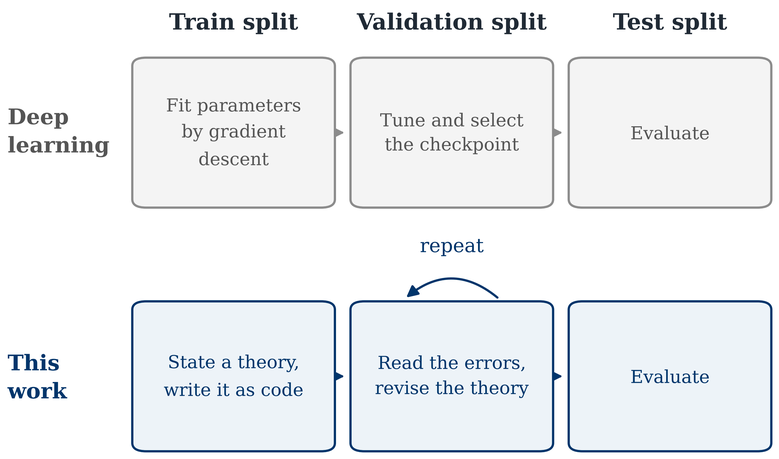
Jeffrey Wei — Yale University
Brian Cheung — University of California, San Francisco

Results & figures

## 1 Replacing the learner

Science deduces the world from first principles, revising its theories where their predictions fail.

We ask whether a reasoning model can form scientific theories, and whether those theories generalize better than deep learning.



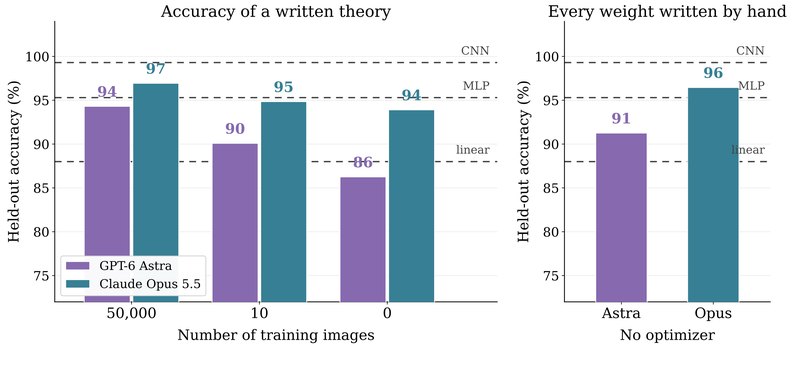

**May** read papers and the web, measure the training data, inspect its own errors

**May not** fit parameters to labels, search automatically over rules, store examples, see the test split

## 2 Reasoning on handwritten digits

MNIST is the test bed. We ask how much data a theory needs, whether a written theory is a better starting point than random initialization, and what the model does when a revision fails.

### How much data a theory needs

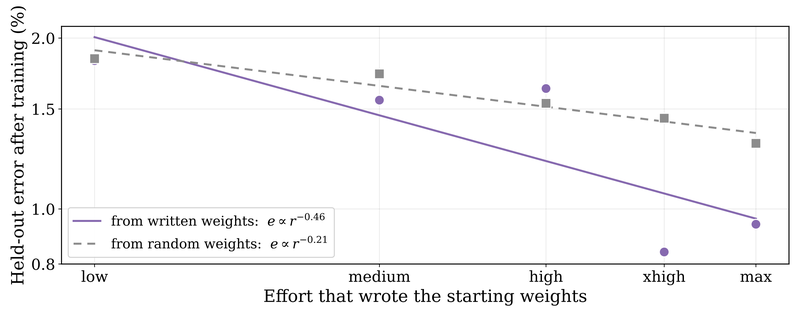



MNIST, 60,000 training and 10,000 test images. Best of five reasoning efforts; dashed lines are trained models.

Given no training images at all, a theory of how digits are written classifies most of them.

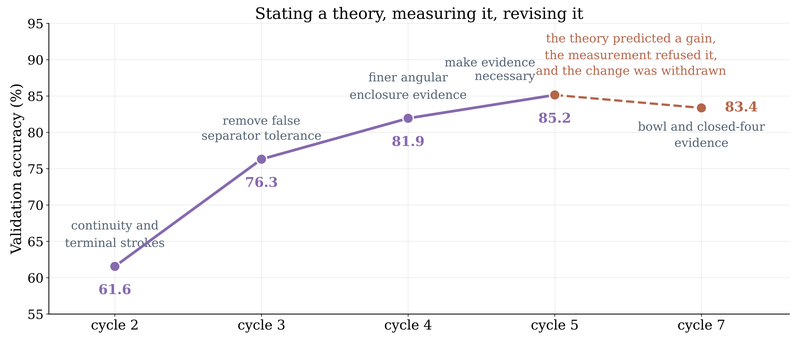
### Reasoning as initialization

Networks initialized from written weights train to lower error, and the gap widens with reasoning effort.




### The deduction–abduction loop

What does a revision look like when the model makes one?

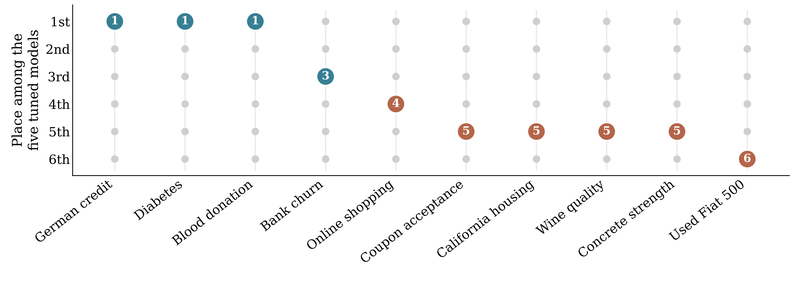



One digits session. Each label is the change the model argued for before measuring.

It abduces a cause, deduces a change, and withdraws the change when the measurement refuses it.

## 3 The same protocol, ten fields

Does the result survive outside images, in fields that differ in how much theory they already have?

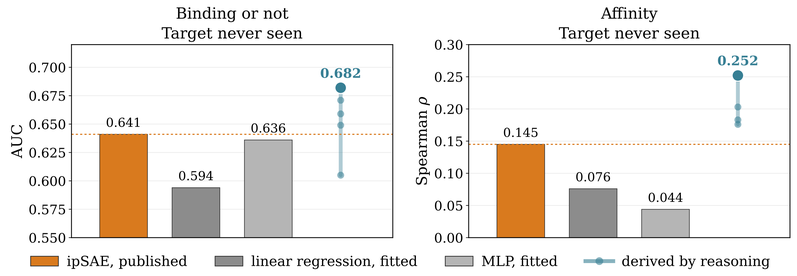



Ten tasks from TabArena, ranked against five tuned models on the same rows.

Wins on small clinical tables, trails boosted trees on feature interactions.

## 4 Deriving a scientific score

Can a model derive a binding score from physics rather than fit one? Each session receives a structure predictor's error matrix and nothing else.




704 designed binders, holding out a whole target. Detection on the left, affinity on the right.

A rule derived from physics beats the published score on a target it has never seen, and beats both fitted models.

### The theory it stated

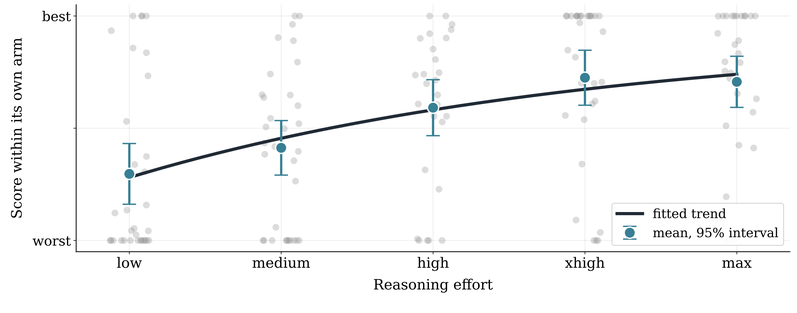
Treat the bound complex as a harmonic well. Equipartition makes the interface stiffness the inverse square of its fluctuation, and a stiffer well is a deeper one. Dividing by the binder's own internal error removes the predictor, and the units with it. Association confines six rigid-body degrees of freedom, each worth half the thermal energy times the log of the stiffness.

$$\kappa = \left(\frac{\sigma_0}{\sigma_{\text{iface}}}\right)^2, \qquad \mathrm{p}K_{\mathrm{d}} = \mathrm{p}K_{\mathrm{d,ref}} + 3\log_{10}\kappa.$$

The three is six halved. Nothing is fitted.

## 5 Scaling reasoning effort

Does more reasoning produce better theories?




## 6 So how far can reasoning go?

We hypothesize that as language model reasoning improves, we will develop stronger first principles in the theory-poor, data-rich domains where deep learning currently dominates.